

# Introduction to Unix Sockets

---

- A socket is a type of descriptor that defines a bi-directional endpoint of communication
- Sockets are used as a basic building block for interprocess communication
- Associated with a socket is a data structure that includes a send buffer and a receive buffer
- A socket is created using the `socket()` system call, which takes as parameters:
  - ☐ protocol family
  - ☐ type of communication
  - ☐ specific protocol (often implicit)
- A socket is created without a name (address); a name is later *bound* to the socket

# Ports

---

- A “port” is an abstraction, provided by both UDP and TCP, that defines a destination endpoints. Think of a port as a mailbox number. Some ports are reserved for specific services and are “published.” (see `/etc/services`)
  - ☐ a socket (an OS-specific entity) can be mapped to a port (a transport protocol entity)
  - ☐ packets arriving for a particular port are queued (in the OS) until a process extracts them
  - ☐ processes waiting at a port are typically blocked until packets arrive
- To communicate with a remote process, the sender must know the internet address of the destination machine as well as the port number that the process is waiting on.

# Socket Naming

---

- Every socket must be *bound* with a name (or address) before it can be referenced
- Different protocols associated with a socket may have different naming structures
- The name space defined by a protocol family is called a *domain*
- The most commonly used families are:
  - Unix Domain (AF\_UNIX)
    - UNIX system internal protocols
    - interprocess communication within same host and file system
    - socket name is a file pathname
  - Internet Domain (AF\_INET)
    - interprocess communication among different hosts
    - socket name includes internet address and port
    - some port numbers are reserved for system use (see `/etc/services`)
    - the name of a socket may be obtained using the system call `getsockname()`

# Primary Socket Types

---

- The socket has the indicated *type*, which specifies the semantics of communication. Currently defined types are:
  - datagram socket (SOCK\_DGRAM)
    - ☐ unreliable datagram communication
    - ☐ record boundary (fixed maximum) is preserved
    - ☐ implemented using the UDP protocol
  - stream socket (SOCK\_STREAM)
    - ☐ reliable virtual circuit communication
    - ☐ “record” boundary is *not* preserved
    - ☐ implemented using the TCP protocol
  - raw socket (SOCK\_RAW)
    - ☐ direct interface to the IP protocol
    - ☐ used in new protocol development
    - ☐ super-user only

# Using Sockets with UDP

---

- Specify use of datagrams upon socket creation.
- If only one-way communication, the sending process need not bind a name to its socket.
- The receiving process *must* bind a name to its socket. The sending process references the name of the receiver's socket.
- Procedures:
  - ☐ socket creation: *socket(domain, type, protocol)*
  - ☐ name (address) binding: *bind(socket, name)*
  - ☐ send or receive through the socket: (several primitives available; see man pages)
  - ☐ close the socket: *close(socket)*
- Sending and receiving datagrams
  - ☐ *sendto()*: used for unconnected sockets
  - ☐ *read()*: primitive routine for sockets and other entities
  - ☐ *recvfrom()*: also retrieves name of sending socket

# Datagram Socket Example - Receiver

---

```
// recv-dgram.cc -- datagram receive code
// After creating a port and binding a name to it, this program prints the port
// number, which is to be used by the sending side. In a loop, the program waits
// for a line of text then sends reply, rot13 encoded.

<include files omitted>

#define BUFLLEN 356

int rot13 ( char *inbuf, char *outbuf ) ;

void main ( )
{
    int      sk ;                // socket descriptor
    sockaddr_in remote ;        // socket address for remote
    sockaddr_in local ;         // socket address for us
    char      buf[BUFLLEN] ;     // buffer from remote
    char      retbuf[BUFLLEN] ;  // buffer to remote
    int       rlen = sizeof(remote) ; // length of remote address
    int       len = sizeof(local) ; // length of local address
    int       moredata = 1 ;     // keep processing or quit
    int       mesglen ;          // actual length of message

    // create the socket
    sk = socket(AF_INET,SOCK_DGRAM,0) ;

    // set up the socket
    local.sin_family = AF_INET ; // internet family
    local.sin_addr.s_addr = INADDR_ANY ; // wild card machine address
    local.sin_port = 0 ;         // let system choose the port

    // bind the name (address) to a port
    bind(sk,(struct sockaddr *)&local,sizeof(local)) ;

    // get the port name and print it out
    getsockname(sk,(struct sockaddr *)&local,&len) ;
    cout << "socket has port " << local.sin_port << "\n" ;
    // cout << "socket has addr " << local.sin_addr.s_addr << "\n" ;
}
```

# Datagram Socket Example (cont.)

---

```
while( moredata ) {
    // wait for a message and print it
    mesglen = recvfrom(sk,buf,BUFLEN,0,(struct sockaddr *)&remote, &rlen) ;
    buf[mesglen] = '\0' ;
    cout << buf << "\n" ;
    moredata = rot13(buf,retbuf) ;

    if( moredata ) {
        // send a reply, using the address given in remote
        sendto(sk,retbuf,strlen(retbuf),0,(struct sockaddr *)&remote, sizeof(remote));
    }
}

/* close the socket */
close(sk);
}

/*
 * Encode message using rot13 scheme.
 */

int rot13 ( char *inbuf, char *outbuf ) {
    int idx ;

    if( inbuf[0]=='.' ) return 0 ;

    idx=0 ;

    while( inbuf[idx]!='\0' ) {
        if( isalpha(inbuf[idx]) ) {
            if( (inbuf[idx]&31)<=13 )
                outbuf[idx] = inbuf[idx]+13 ;
            else
                outbuf[idx] = inbuf[idx]-13 ;
        } else
            outbuf[idx] = inbuf[idx] ;
        idx++ ;
    }
    outbuf[idx] = '\0';
    return 1 ;
}
```

# Datagram Socket Example - Sender

---

```
// send-dgram.cc -- datagram sending code
// This program sends a message in a datagram, waits for, and prints a reply. The
// destination machine name and destination port number are command line arguments.

#include files deleted>

#define MSG1 "Have you heard about the new corduroy pillows?\nThey are making headlines!"
#define MSG2 "."
#define BUFLen 356

void main ( int argc, char *argv[] )
{
    int          sk ;           // socket descriptor
    sockaddr_in  remote ;       // socket address for remote side
    char         buf[BUFLen] ;   // buffer for response from remote
    hostent      *hp ;          // address of remote host
    int          mesglen ;       // actual length of the message

    // create the socket
    sk = socket(AF_INET,SOCK_DGRAM,0) ;
    // designate the addressing family
    remote.sin_family = AF_INET ;
    // get the address of the remote host and store
    hp = gethostbyname(argv[1]) ;
    memcpy(&remote.sin_addr, hp->h_addr, hp->h_length) ;
    remote.sin_port = atoi(argv[2]) ;

    // send the message to the other side
    sendto(sk,MSG1,strlen(MSG1),0,(struct sockaddr *)&remote, sizeof(remote)) ;

    // wait for a response and print it
    mesglen = read(sk,buf,BUFLen) ;
    buf[mesglen] = '\0';
    cout << buf << "\n" ;

    // send message telling it to shut down
    sendto(sk,MSG2,strlen(MSG2),0,(struct sockaddr *)&remote, sizeof(remote)) ;

    // close the socket and exit
    close(sk);
}
```



# Structures Defined in System Header Files

---

```
struct sockaddr {
    unsigned short    sa_family;    // address family, AF_XXX
    char              sa_data[14];  // 14 bytes of protocol address
};

// IPv4 AF_INET sockets:

struct sockaddr_in {
    short             sin_family;    // e.g. AF_INET, AF_INET6
    unsigned short    sin_port;      // e.g. htons(3490)
    struct in_addr     sin_addr;     // see struct in_addr, below
    char              sin_zero[8];   // zero this if you want to
};

struct in_addr {
    unsigned long s_addr;            // load with inet_pton()
};

struct hostent
{
    char *h_name; /* Official name of host. */
    char **h_aliases; /* Alias list. */
    int h_addrtype; /* Host address type. */
    int h_length; /* Length of address. */
    char **h_addr_list; /* List of addresses from name server. */
    #define h_addr h_addr_list[0] /* Address, for backward compatibility. */
};
```

# Unix Stream Sockets

---

- Also created with *socket()* system call
- Socket type: SOCK\_STREAM
  - ☐ reliable virtual circuit communication
  - ☐ “record” boundary is *not* preserved
  - ☐ When used in AF\_INET domain, implemented using the TCP protocol
- *socketpair()* - create a pair of connected sockets
  - ☐ generalization of *pipes*
  - ☐ only supports the AF\_UNIX domain
  - ☐ same machine
- Timeouts and broken connections: If data is not transmitted within a reasonable length of time, then the connection is broken and subsequent calls will fail with ETIMEDOUT.
- Use *bind()* system call for both datagram and stream sockets
  - ☐ socket descriptor
  - ☐ name (family, host addr, port)
  - ☐ length of name
  - ☐ often let the OS choose the port

# Stream Socket – Passive Side

---

- Await (listen for) connections
- Use *listen()* system call to prepare operating system for connection requests
- Parameters:
  - ☐ socket descriptor
  - ☐ backlog: defines the maximum length the queue of pending connections (usually 5)

## Passive Side (cont.)

---

- Use the *accept()* system call
- Parameters
  - ☐ socket descriptor
  - ☐ name (family, host addr, port)
  - ☐ length of name
- Extracts the first connection on the queue of pending connections, creates a new socket, and allocates a new file descriptor for the socket.
- Normally blocks, but this can be turned off
- Returns descriptor of new socket.
- Why does *accept()* work this way?

# Stream Socket - Active Side

---

- Initiate a connection on a socket
- Use the *connect()* system call
- Parameters:
  - ☐ socket descriptor
  - ☐ name (family, host addr, port)
  - ☐ length of name
- For stream sockets, attempts to make a connection to the socket named in the call.
- Errors
  - ☐ ETIMEOUT
  - ☐ ECONNREFUSED
  - ☐ ENETDOWN or EHOSTDOWN
  - ☐ ENETUNREACH or EHOSTREACH

# Transferring Data

---

- Transmitting
  - ☐ *write()* as for files
  - ☐ *send()* uses flags to specify such requests as out-of-band transmission (MSG\_OOB)
- Receiving
  - ☐ *read()* as for files
  - ☐ *recv()* uses flags to specify such requests as examining data without reading it (MSG\_PEEK)

# Stream Example - Passive Side

---

```
// recv-stream.cc -- passive side stream socket example
void main ( )
{
    int          sk, sk2 ;                // socket descriptors
    sockaddr_in  remote ;                // socket address for remote
    sockaddr_in  local ;                 // socket address for us
    char         buf[BUFLEN] ;           // buffer from remote
    char         retbuf[BUFLEN] ;        // buffer to remote
    int          rlen = sizeof(remote) ; // length of remote address
    int          len = sizeof(local) ;   // length of local address
    int          moredata = 1 ;          // keep processing or quit
    int          mesglen ;                // actual length of message

    // create the socket
    sk = socket(AF_INET,SOCK_STREAM,0) ;
    // set up the socket
    local.sin_family = AF_INET ;         // internet family
    local.sin_addr.s_addr = INADDR_ANY ; // wild card machine address
    local.sin_port = 0 ;                 // let system choose the port
    // bind the name (address) to a port
    bind(sk,(struct sockaddr *)&local,sizeof(local)) ;
    // get the port name and print it out
    getsockname(sk,(struct sockaddr *)&local,&len) ;
    cout << "socket has port " << local.sin_port << "\n" ;
    // tell OS to queue (up to 1) connection requests
    listen(sk, 1);

    // wait for connection request, then close old socket
    sk2 = accept(sk, (struct sockaddr *)0, (int *)0) ;
    close(sk);

    if(sk2 == -1)
        cout << "accept failed!\n" ;
    else { while( moredata ) {
        // wait for a message and print it
        mesglen = read(sk2,buf,BUFLEN);
        buf[mesglen] = '\0' ;
        cout << buf << "\n" ;
        moredata = rot13(buf,retbuf) ;
        if( moredata ) {
            // send a reply
            write(sk2,retbuf,strlen(retbuf));
        } } }
    close(sk2);
    exit(0);
}
```

# Stream Example - Active Side

---

```
// send-stream.cc -- active side stream socket example

#define MSG1 "Have you heard about the new corduroy pillows?\nThey are making headlines!"
#define MSG2 "."
#define BUFLen 356

void main ( int argc, char *argv[] )
{
    int          sk ;           // socket descriptor
    sockadr_in    remote ;      // socket address for remote side
    char          buf[BUFLen] ;  // buffer for response from remote
    hostent       *hp ;         // address of remote host
    int           mesglen ;      // actual length of the message

    // create the socket
    sk = socket(AF_INET,SOCK_STREAM,0) ;
    // designate the addressing family
    remote.sin_family = AF_INET ;
    // get the address of the remote host and store
    hp = gethostbyname(argv[1]) ;
    memcpy(&remote.sin_addr, hp->h_addr, hp->h_length) ;
    // get the port used on the remote side and store
    remote.sin_port = atoi(argv[2]) ;

    // connect to other side
    if(connect(sk, (struct sockadr *)&remote, sizeof(remote)) < 0) {
        cout << "connection error!\n" ;
        close(sk);
        exit(1);
    }

    // send the message to the other side
    write(sk,MSG1,strlen(MSG1));
    // wait for a response and print it
    mesglen = read(sk,buf,BUFLen) ;
    buf[mesglen] = '\0';
    cout << buf << "\n" ;
    // send message telling it to shut down
    write(sk,MSG2,strlen(MSG2));
    // close the socket and exit
    close(sk);
}
```

Why does accept() return a new socket descriptor?



# Server Example - Passive Side

---

```
// recv-server.cc -- handles translation for multiple clients
void main ( )
{

    <variables as in earlier example>

    // create the socket
    sk = socket(AF_INET,SOCK_STREAM,0) ;
    // set up the socket
    local.sin_family = AF_INET ;           // internet family
    local.sin_addr.s_addr = INADDR_ANY ; // wild card machine address
    local.sin_port = 0 ;                   // let system choose the port

    // bind the name (address) to a port
    bind(sk,(struct sockaddr *)&local,sizeof(local)) ;
    // get the port name and print it out
    getsockname(sk,(struct sockaddr *)&local,&len) ;
    cout << "socket has port " << local.sin_port << "\n" ;
    // tell OS to queue (up to 5) connection requests
    listen(sk, 5);

    // we loop forever, taking connections and forking new servers
    for( ; ; ) {
        // wait for connection request
        sk2 = accept(sk, (struct sockaddr *)0, (int *)0) ;
        if(sk2 == -1)    cout << "accept failed!\n" ;
        else {
            if( fork()==0 ) {
                // this is the child process ...
                close(sk) ;           // sk is no longer needed
                while( moredata ) {
                    // wait for a message and print it
                    mesglen = read(sk2,buf,BUFLen);
                    buf[mesglen] = '\0' ;
                    cout << buf << "\n" ;
                    moredata = rot13(buf,retbuf) ;
                    if( moredata ) { // send a reply
                        write(sk2,retbuf,strlen(retbuf));
                    }
                }
                exit(0);
            }
            // this is the parent, so we no longer need sk2 ...
            close(sk2) ;
        }
    }
}
```

# Server Example - Active Side

---

```
// send-server.cc -- active side for server example

#define ENDMSG "."
#define BUFLen 356
void main ( int argc, char *argv[] )
{
    int          sk ;           // socket descriptor
    sockadr_in   remote ;      // socket address for remote side
    char         buf1[BUFLen] ; // buffer for sending to remote
    char         buf2[BUFLen] ; // buffer for response from remote
    hostent      *hp ;         // address of remote host
    int          mesglen ;      // actual length of the message

    // create the socket
    sk = socket(AF_INET,SOCK_STREAM,0) ;
    // designate the addressing family
    remote.sin_family = AF_INET ;
    // get the address of the remote host and store
    hp = gethostbyname(argv[1]) ;
    memcpy(&remote.sin_addr, hp->h_addr, hp->h_length) ;
    // get the port used on the remote side and store
    remote.sin_port = atoi(argv[2]) ;

    // connect to other side
    if(connect(sk, (struct sockadr *)&remote, sizeof(remote)) < 0) {
        cout << "connection error!\n" ;
        close(sk);
        exit(1);
    }
    // loop, reading input and sending to other side, until a single '.' is typed
    cin.getline (buf1, sizeof(buf1));
    while (buf1[0] != '.') {
        // send the message to the other side
        write(sk, buf1, strlen(buf1));
        // wait for a response and print it
        mesglen = read(sk, buf2, BUFLen) ;
        buf2[mesglen] = '\0';
        cout << buf2 << "\n" ;
        // get next line of input
        cin.getline (buf1, sizeof(buf1));
    }
    // send (last) message telling it to shut down
    write(sk, buf1, strlen(buf1));
    // close the socket and exit
    close(sk);
}
```

# Select System Call

---

- Parameters:
  - ☐ number of file descriptors
  - ☐ &readmask
  - ☐ &writemask
  - ☐ &exceptmask
  - ☐ timeout
- Operation
  - ☐ examines the I/O descriptor sets whose addresses are passed as mask parameter
  - ☐ replaces passed mask with map of those ready
  - ☐ passes back the number that are ready
- How can we use this?

# Important Functions

---

Find how to use the following functions:

- `socket(); accept(); bind(); close(); connect(); fork(); listen();`
- `getaddrinfo(); gethostbyaddr(); gethostbyname(); gethostname(); getsockname();`
- `htonl(); htons(); ntohl(); ntohs(); inet_aton(); inet_ntoa(); inet_ntop(); inet_pton();`
- `sendto(); recvfrom(); send(); select();`

Resources:

- 'man' on Unix/Linux machines
- online socket programming guides, e.g., <http://beej.us/guide/bgnet/>
- Unix Network Programming (Volume 1) by W. Richard Stevens)

# Internet Address: data structures & function

- Address: 35.8.10.140
- Name: www.msu.edu
- Representation preferred by programs?
- struct addrinfo {
  - int ai\_flags; // AI\_PASSIVE, AI\_CANONNAME, etc.
  - int ai\_family; // AF\_INET, AF\_INET6, AF\_UNSPEC
  - int ai\_socktype; // SOCK\_STREAM, SOCK\_DGRAM
  - int ai\_protocol; // use 0 for "any"
  - size\_t ai\_addrlen; // size of ai\_addr in bytes
  - struct sockaddr \*ai\_addr; // struct sockaddr\_in or \_in6
  - char \*ai\_canonname; // full canonical hostname
  - struct addrinfo \*ai\_next; // linked list, next node
  - };
- Function getaddrinfo( )

```
struct sockaddr {  
    unsigned short    sa_family; // address family, AF_XXX  
    char             sa_data[14]; // 14 bytes of protocol address  
};
```

```
struct sockaddr_in {  
    short int         sin_family; // Address family, AF_INET  
    unsigned short int sin_port;  // Port number  
    struct in_addr     sin_addr;  // Internet address  
    unsigned char      sin_zero[8]; // Same size as struct sockaddr  
};
```

// (IPv4 only--see struct in6\_addr for IPv6)

// Internet address (a structure for historical reasons)

```
struct in_addr { uint32_t s_addr; // that's a 32-bit int (4 bytes) };
```

## **struct hostent {**

- `char *h_name;` //This is the "official" name of the host.
  - `char **h_aliases;` //These are alternative names for the host,  
represented as a null-terminated vector of strings.
  - `int h_addrtype;` //This is the host address type; in practice, its  
value is always either `AF_INET` or `AF_INET6`
  - `int h_length;` //This is the length, in bytes, of each address.
  - `char **h_addr_list;` //This is the vector of addresses for the host.
  - `char *h_addr;` //This is a synonym for `h_addr_list[0]`;
- }



# Important Function.

Find how to use the following functions:

- `socket(); accept(); bind(); close(); connect(); fork(); listen();`
- `getaddrinfo(); gethostbyaddr(); gethostbyname(); gethostname(); getsockname();`
- `htonl(); htons(); ntohl(); ntohs(); inet_aton(); inet_ntoa(); inet_ntop(); inet_pton();`
- `sendto(); recvfrom(); send(); recv(); select();`

Resources:

- 'man' on Unix/Linux machines
- online socket programming guides, e.g., <http://beej.us/guide/bgnet/>
- Unix Network Programming (Volume 1) by W. Richard Stevens, online at: <http://proquestcombo.safaribooksonline.com.proxy2.cl.msu.edu/0-13-141155-1>